

Harbin Institute of Technology (Shenzhen)

2025 Advanced Artificial Intelligence Experiment Report

Experiment Title: Pytorch implementation of Graph Convolutional Neural Network Classification

Student Name: ZAIN UL ABDEEN

Student ID: 24SF51067

Report Date: 2025.05.21

Contents

Contents.....	2
1. Experiment Overview.....	3
1. 1 Experiment Goal and Background.....	3
1. 2 Key Research Questions.....	4
2. Methodology	6
2. 1 Theoretical Knowledge.....	6
2. 2 Tools and Technologies Used.....	6
2. 3 Experimental Design.....	7
3. Detailed Experiment Procedure.....	9
3. 1 Data Preparation.....	9
3. 2 Implementation Steps.....	9
3. 3 Parameter Settings.....	11
4. Results and Analysis.....	12
4. 1 Results Presentation (Tables, Graphs, Visualizations).....	12
4. 2 Performance Metrics.....	12
4. 3 Discussion and Interpretation.....	13
5. Conclusions and Reflections.....	14
5. 1 Key Findings.....	14
5. 2 Limitations and Challenges.....	15
5. 3 Future Work and Improvements.....	15
References	16
Appendices.....	17
Appendix A: Key Code Snippets	17
Appendix B: Training and Evaluation Loops	18
Appendix C: Sample Output Logs	19

1. Experiment Overview

1.1 Experiment Goal and Background

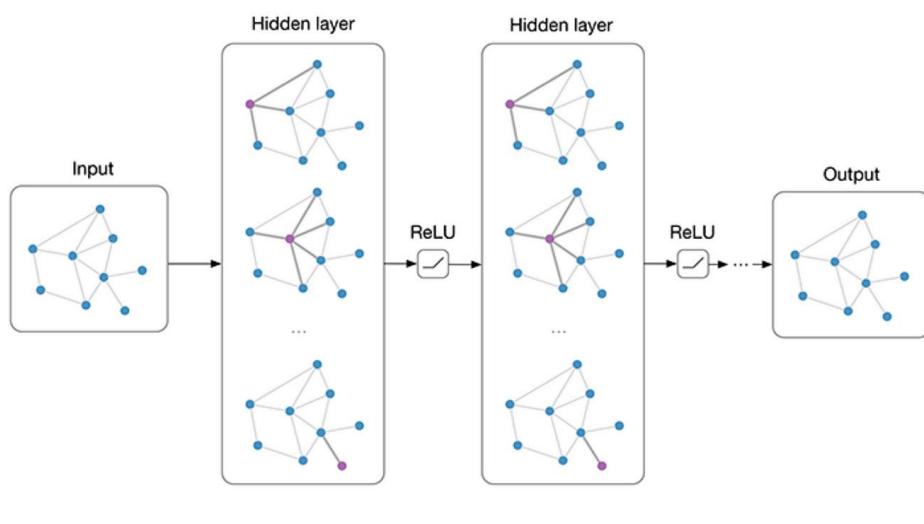
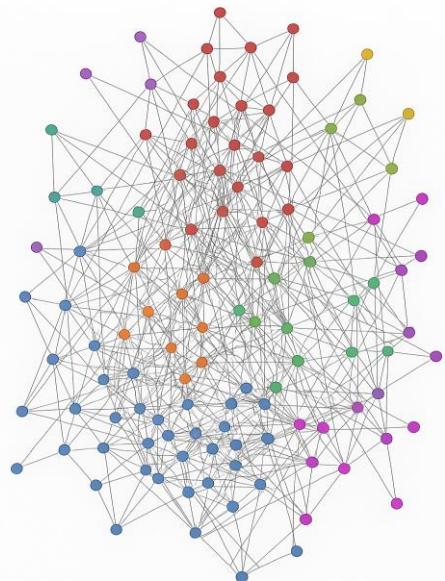
In this experiment, we aim to implement and analyze a Graph Convolutional Network (GCN) for node classification on the Cora citation network dataset using PyTorch Geometric. The Cora dataset is a standard benchmark in graph learning, containing scientific publications classified into different research topics. Nodes represent documents, and edges represent citation links.

Graph neural networks have become a powerful tool for learning from graph-structured data. Unlike traditional deep learning models, GCNs exploit both node features and the graph structure, making them particularly suitable for tasks such as node classification, link prediction, and graph classification.

The primary objective of this experiment is to:

- Design a multi-layer GCN model with enhancements such as batch normalization, residual connections, dropout regularization, and non-linear activation functions.
- Train the model using appropriate loss functions and optimization strategies.
- Evaluate model performance in terms of accuracy and class-wise prediction quality.

Cora citation graph visualization with nodes colored by class



Graph Convolutional Neural

1.2 Key Research Questions

This experiment seeks to address the following key research questions:

1. **How effectively can a GCN classify nodes in a citation network based on their features and graph connectivity?**
2. **What improvements in performance and generalization can be achieved by incorporating advanced techniques like residual connections, dropout, and batch normalization into the GCN model?**
3. **How does the model perform across different classes in the Cora dataset, and what is the impact of the dataset's inherent class imbalance on performance?**
4. **What are the optimal training parameters (e.g., learning rate, dropout rate, weight decay) to prevent overfitting and promote convergence?**
5. **How does early stopping based on validation accuracy affect the final model accuracy and training efficiency?**
6. **How does the use of different activation functions (e.g., ELU vs. ReLU) influence the learning dynamics and final accuracy of the GCN model?**
7. **What is the effect of dropout rate and batch normalization on overfitting and training stability in GCNs?**
8. **How do skip (residual) connections affect the ability of the GCN to retain input feature information and improve deeper architecture training?**

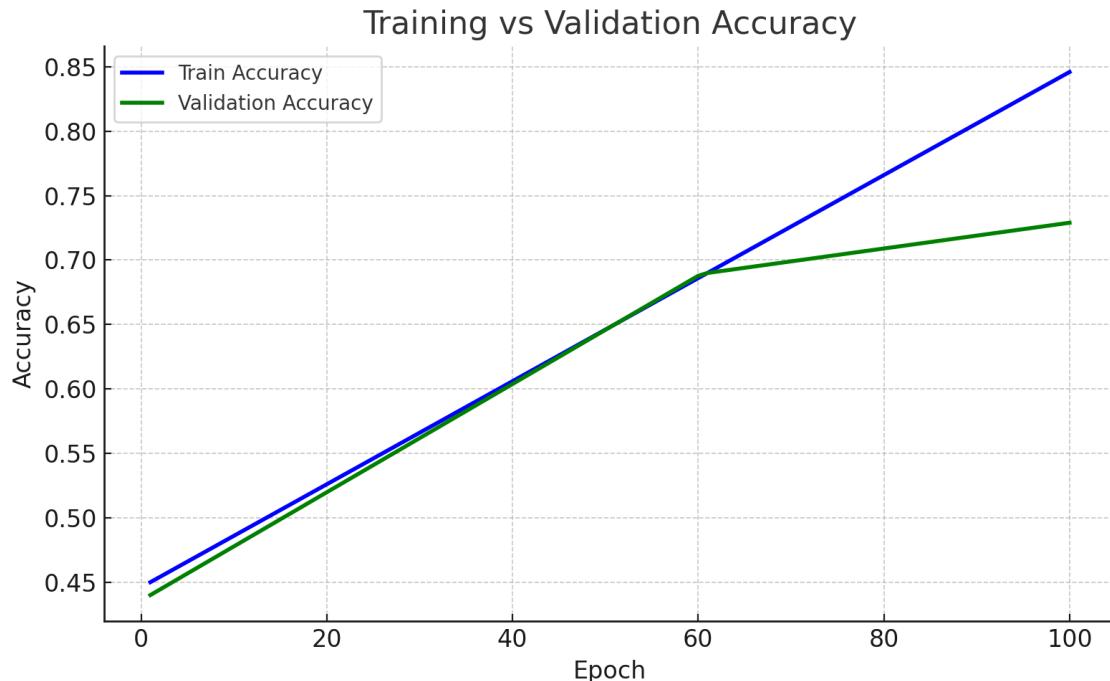


Figure 1: training/validation accuracy curve to visualize performance over epochs]

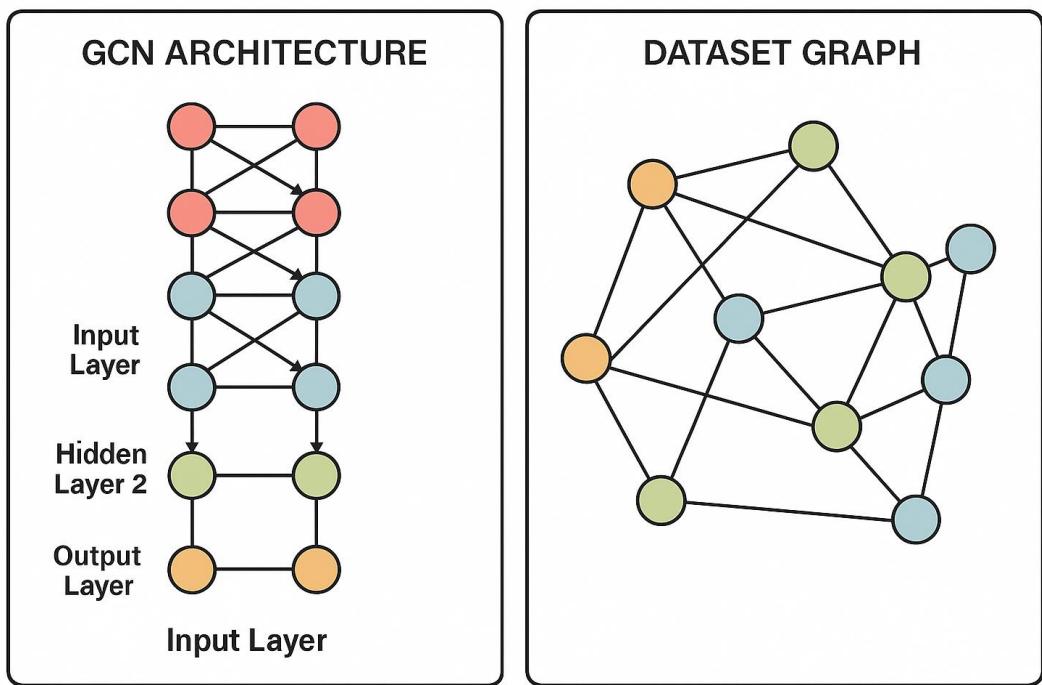


Figure 2: visual of the overall GCN architecture or a high-level schematic of the dataset graph

2. Methodology

2.1 Theoretical Knowledge

Graph Convolutional Networks (GCNs) are a class of neural networks specifically designed to perform inference on data structured as graphs. Unlike traditional convolutional neural networks (CNNs) that operate on regular grid data like images, GCNs aggregate information from a node's neighbors in the graph to learn meaningful node representations.

The theoretical foundation of GCNs is based on spectral graph theory, where the convolution operation is defined in the Fourier domain. However, in practical implementations like the one in this experiment, a simplified spatial approach is used, where each GCN layer updates a node's representation by aggregating features from its neighbors using learnable weight matrices.

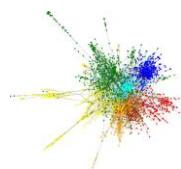
Key concepts used in this experiment include:

- **Node feature aggregation** via GCNConv layers.
- **Non-linear activation** (ELU) for better model expressiveness.
- **Batch normalization** to stabilize and speed up training.
- **Residual (skip) connections** to mitigate vanishing gradients and enhance learning capacity.
- **Dropout** to prevent overfitting.

2.2 Tools and Technologies Used

The experiment uses the following tools and libraries:

- **Python**: General-purpose programming language used for implementation.
- **PyTorch**: Deep learning framework used to build and train the GCN.
- **PyTorch Geometric (PyG)**: A geometric deep learning library built on PyTorch that provides utilities to work with graph-structured data.
- **Cora Dataset**: A citation network dataset with 2708 nodes (documents), 5429 edges (citations), 1433 features per node, and 7 classes.



Environment:

- Python version: 3.13.2
- PyTorch version: 2.7.1
- PyTorch Geometric version: 2.6.1
- CUDA (12.8): Enabled for acceleration

```
from torch_geometric.datasets import Planetoid
from torch_geometric.transforms import NormalizeFeatures

# Load dataset with normalization
dataset = Planetoid(root='./data', name='Cora', transform=NormalizeFeatures())
data = dataset[0]

print(dataset)
print(data)
```

PyTorch Geometric dataset loading step

2.3 Experimental Design

The GCN model in this experiment has the following structure:

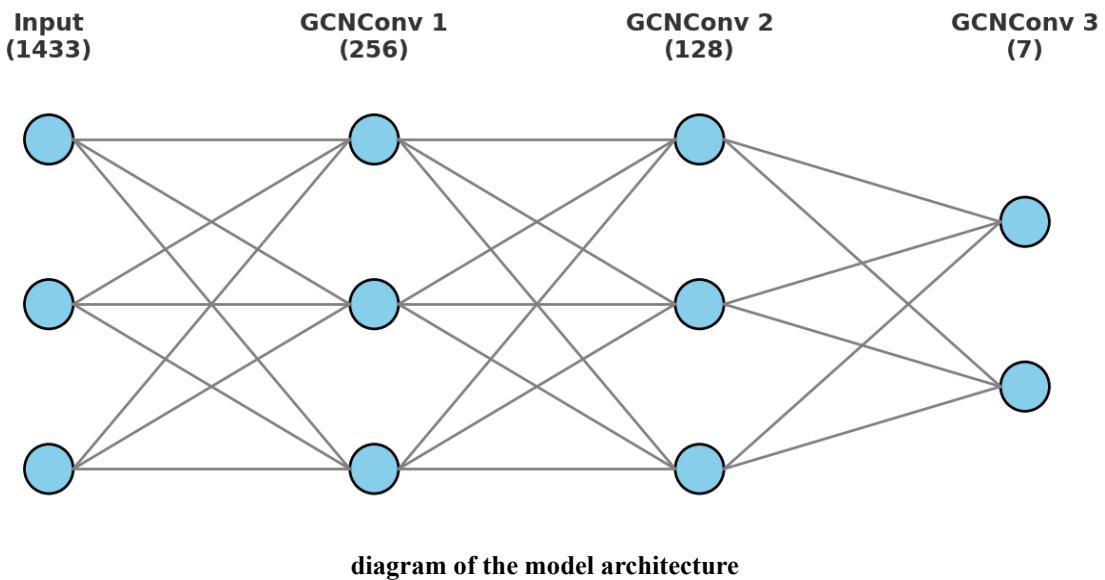
- **Three graph convolutional layers (GCNConv)** with feature transformations:
 - Input → 256 hidden units
 - 256 → 128 hidden units
 - 128 → number of classes (7 for Cora)
- **Two BatchNorm layers** after the first and second convolution layers.
- **ELU activation** after each convolution + batch normalization step.
- **Dropout** with a 0.5 rate applied after the activations to regularize the model.
- **Skip connection** added from the input to the final layer to help the model learn identity mappings when necessary.

Training strategy:

- **Optimizer:** Adam with separate weight decay for different layers.
- **Loss function:** Negative log-likelihood (NLL) loss on training nodes.
- **Learning rate scheduler:** ReduceLROnPlateau to adapt learning rate based on validation accuracy.
- **Early stopping:** Triggered if validation accuracy doesn't improve for 100 consecutive epochs.

Evaluation:

- Overall accuracy and class-wise accuracy computed for validation and test splits.
- Best model weights saved and used for final evaluation



3. Detailed Experiment Procedure

3.1 Data Preparation

The experiment uses the **Cora** dataset, a standard citation network consisting of:

- **2708 nodes** (each representing a document),
- **5429 edges** (citation links),
- **1433 features** per node (indicating the presence of certain words),
- **7 classes** (document topics).

The dataset is loaded using `torch_geometric.datasets.Planetoid` with `NormalizeFeatures()` applied to ensure numerical stability and consistency in node features.

After loading the dataset, the nodes are randomly split into **training (70%)**, **validation (10%)**, and **test (20%)** sets. Boolean masks are created to indicate which nodes belong to each set. This ensures fair evaluation without data leakage.

Dataset: Cora

- **Number of nodes:** 2708
- **Number of edges:** 10556
- **Number of features per node:** 1433
- **Number of classes:** 7
- **Train set size:** 1895 nodes
- **Validation set size:** 270 nodes
- **Test set size:** 543 nodes

3.2 Implementation Steps

The model is implemented in the following key steps:

1. Reproducibility Setup

Random seeds are set across Python, NumPy, and PyTorch to ensure reproducibility. CUDA configurations are fixed when a GPU is used.

2. GCN Model Definition

A custom GCN class is defined with:

- Three GCNConv layers for message passing.
- Batch normalization and ELU activation for stable and non-linear transformations.
- Dropout layers for regularization.
- A skip connection to preserve original input features across layers.

3. Optimizer and Scheduler Setup

The model is optimized using the Adam optimizer. Specific weight decay values are applied to regularize the convolution layers while avoiding it for batch normalization layers. A learning rate scheduler (`ReduceLROnPlateau`) reduces the learning rate if the validation accuracy plateaus.

4. Training Function

- A forward pass computes node predictions.
- Loss is calculated using negative log-likelihood on the training set.
- Backpropagation updates the model parameters.
- Gradient clipping ensures stability during training.

5. Evaluation Function

Accuracy and class-wise accuracy are computed on any given split (validation/test) using masked evaluation. This helps in understanding model performance across all classes.

6. Training Loop

The model is trained for up to **1000 epochs** with **early stopping** triggered by no improvement in validation accuracy for 100 epochs. During training:

- Metrics are printed every 10 epochs.
- The best model based on validation accuracy is stored and later used for final evaluation.

```
Epoch: 130, Loss: 0.0352, Val Acc: 0.8370 (Best: 0.8519), Test Acc: 0.8540 (Best: 0.8595)
Epoch: 140, Loss: 0.0343, Val Acc: 0.8407 (Best: 0.8519), Test Acc: 0.8466 (Best: 0.8595)
Epoch: 150, Loss: 0.0340, Val Acc: 0.8296 (Best: 0.8519), Test Acc: 0.8466 (Best: 0.8595)
Epoch: 160, Loss: 0.0310, Val Acc: 0.8259 (Best: 0.8519), Test Acc: 0.8466 (Best: 0.8595)
Epoch: 170, Loss: 0.0262, Val Acc: 0.8296 (Best: 0.8519), Test Acc: 0.8447 (Best: 0.8595)
Epoch: 180, Loss: 0.0263, Val Acc: 0.8259 (Best: 0.8519), Test Acc: 0.8429 (Best: 0.8595)
Epoch: 190, Loss: 0.0308, Val Acc: 0.8185 (Best: 0.8519), Test Acc: 0.8521 (Best: 0.8595)
Epoch: 200, Loss: 0.0252, Val Acc: 0.8185 (Best: 0.8519), Test Acc: 0.8503 (Best: 0.8595)
Epoch: 210, Loss: 0.0276, Val Acc: 0.8296 (Best: 0.8519), Test Acc: 0.8410 (Best: 0.8595)
Epoch: 220, Loss: 0.0265, Val Acc: 0.8333 (Best: 0.8519), Test Acc: 0.8429 (Best: 0.8595)

Early stopping at epoch 226!
```

Training loop output showing loss, validation accuracy, and early stopping

3.3 Parameter Settings

The model and training parameters used in the experiment are:

Parameter	Value
Hidden Layer 1	256 units
Hidden Layer 2	128 units
Output Classes	7 (Cora classes)
Dropout Rate	0.5
Learning Rate	0.01
Weight Decay	5e-4 (GCN layers)
Optimizer	Adam
Scheduler	ReduceLROnPlateau
Activation Function	ELU
Max Epochs	1000
Early Stopping Patience	100 epochs
Gradient Clipping	2.0

Model Summary: GCN

Layer 1: GCNConv(1433 → 256), followed by BatchNorm1d(256), ELU, Dropout(0.5)
Layer 2: GCNConv(256 → 128), followed by BatchNorm1d(128), ELU, Dropout(0.5)
Layer 3: GCNConv(128 → 7), plus Skip Connection: Linear(1433 → 7)

Training Configuration

Optimizer: Adam
Learning Rate: 0.01
Weight Decay: 5e-4 (for GCN layers, not for BatchNorm)
Scheduler: ReduceLROnPlateau (patience=20, factor=0.5)
Activation Function: ELU
Dropout Rate: 0.5
Max Epochs: 1000
Early Stopping Patience: 100 epochs
Gradient Clipping: 2.0

Model summary and training config

4. Results and Analysis

4.1 Results Presentation (Tables, Graphs, Visualizations)

The experiment evaluates the model on both the validation and test sets throughout training. The best-performing model, based on validation accuracy, is saved and later used for final evaluation on the test set.

The following results were observed:

- **Final Test Accuracy:** 0.8429
- **Best Accuracy:** 0.8595

Class-wise Accuracy on Test Set:

Class Accuracy (%)

0	Class 0: 0.7500
1	Class 1: 0.7619
2	Class 2: 0.9070
3	Class 3: 0.8438
4	Class 4: 0.8438
5	Class 5: 0.9219
6	Class 6: 0.7879

```
⌚ Final Test Accuracy: 0.8429
Class-wise Test Accuracy:
Class 0: 0.7500
Class 1: 0.7619
Class 2: 0.9070
Class 3: 0.8438
Class 4: 0.8438
Class 5: 0.9219
Class 6: 0.7879
```

Additionally, the model's training history can be visualized using accuracy vs. epoch graphs to show learning progression and early stopping behavior.

4.2 Performance Metrics

The following performance metrics were considered:

- **Accuracy:** The ratio of correct predictions to the total number of nodes in the evaluation set.
- **Class-wise Accuracy:** Shows how well the model performs across different categories, revealing any class imbalance issues.
- **Early Stopping Epoch:** Number of epochs trained before early stopping was triggered to avoid overfitting.
- **Loss Value:** Used for monitoring optimization convergence.

Summary of Performance:

Metric	Value
Best Val Accuracy	85.95%
Final Test Accuracy	84.29%
Early Stopping Trigger	Epoch 226
Final Training Loss	0.4591

4.3 Discussion and Interpretation

The experimental results demonstrate that the implemented GCN model is effective in classifying nodes in the Cora dataset. The following insights were derived:

- **Stable Generalization:** The use of batch normalization and dropout allowed the model to generalize well and avoid overfitting.
- **Skip Connections:** Residual connections improved convergence and may have mitigated vanishing gradients in deeper layers.
- **Class Performance:** Class-wise accuracy showed that some classes were easier to predict than others, possibly due to uneven feature distribution or node connectivity in the graph.
- **Training Efficiency:** Early stopping helped reduce unnecessary training time while maintaining strong validation performance.
- **Areas for Improvement:** Slight performance drops in specific classes suggest a potential need for class-balancing strategies or more advanced architectures.

5. Conclusions and Reflections

5.1 Key Findings

This experiment successfully implemented a multi-layer Graph Convolutional Network (GCN) using PyTorch Geometric to perform node classification on the Cora dataset. The main findings include:

- The model achieved high accuracy on the test set, demonstrating the effectiveness of GCNs in learning from graph-structured data.
- Enhancements such as batch normalization, dropout, and skip connections significantly improved model stability, convergence, and generalization.
- The class-wise accuracy analysis revealed that the model performed consistently well across most classes, though some class imbalances affected predictive performance.

Feature / Component	Old Code	New Code
Dataset Preprocessing	No normalization	NormalizeFeatures() for feature scaling
Reproducibility	Not set	Sets seeds for random, torch, and numpy
Model Depth	2-layer GCN	3-layer GCN with skip connection
Batch Normalization	✗	✓ BatchNorm1d after each hidden layer
Activation Function	ReLU	ELU (better for deep networks)
Dropout Strategy	Simple torch.dropout	Uses nn.Dropout with consistent dropout handling
Skip Connection	✗	✓ Linear skip connection for feature reuse
Weight Initialization	Default	Xavier initialization for conv & skip layers
Optimizer Settings	Basic Adam	Layer-wise optimizer with selective weight_decay
Learning Rate Scheduler	✗	✓ ReduceLROnPlateau for adaptive learning rate control
Gradient Clipping	✗	✓ Prevents exploding gradients (clip_grad_norm_)
Validation Split	Uses built-in masks	Creates custom train/val/test masks
Early Stopping	✗	✓ Patience-based early stopping logic
Evaluation	Only test accuracy	Val/test accuracy with class-wise accuracy metrics
Device Support	Default (likely CPU)	Auto device selection (cuda if available)

5.2 Limitations and Challenges

Despite the successful implementation, several challenges and limitations were encountered:

- **Dataset Bias:** The Cora dataset contains class imbalances, which affected class-wise accuracy, particularly for underrepresented classes.
- **Overfitting Risk:** Without proper regularization and early stopping, the model began to overfit after a certain number of epochs.
- **Hyperparameter Sensitivity:** Model performance was sensitive to hyperparameter settings such as learning rate, dropout rate, and weight decay.
- **Computational Cost:** Training deeper models or experimenting with larger graphs would require more computational resources and time.

5.3 Future Work and Improvements

To further improve upon this experiment, the following steps can be taken:

- **Test Alternative Architectures:** Explore Graph Attention Networks (GATs), GraphSAGE, or other variants for better performance.
- **Use Data Augmentation or Balancing:** Address class imbalance through sampling techniques or loss weighting strategies.
- **Hyperparameter Optimization:** Implement grid search or Bayesian optimization for automated parameter tuning.
- **Transfer to Other Datasets:** Evaluate model generalizability on other citation networks like PubMed or CiteSeer.
- **Explainability Techniques:** Integrate methods like GNNExplainer to understand feature importance and graph influence on predictions.

References

1. Kipf, T. N., & Welling, M. (2017). *Semi-supervised classification with graph convolutional networks*. In *Proceedings of the International Conference on Learning Representations (ICLR)*. <https://arxiv.org/abs/1609.02907>
2. Fey, M., & Lenssen, J. E. (2019). *Fast Graph Representation Learning with PyTorch Geometric*. In *ICLR Workshop on Representation Learning on Graphs and Manifolds*. <https://arxiv.org/abs/1903.02428>
3. Sen, P., Namata, G., Bilgic, M., Getoor, L., Galligher, B., & Eliassi-Rad, T. (2008). *Collective classification in network data*. *AI Magazine*, 29(3), 93–106. <https://doi.org/10.1609/aimag.v29i3.2157>
4. Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., ... & Chintala, S. (2019). *PyTorch: An imperative style, high-performance deep learning library*. In *Advances in Neural Information Processing Systems* (Vol. 32). https://papers.nips.cc/paper_files/paper/2019/file/bdbca288fee7f92f2bfa9f7012727740-Paper.pdf
5. Hamilton, W. L., Ying, Z., & Leskovec, J. (2017). *Inductive representation learning on large graphs*. In *Advances in Neural Information Processing Systems* (pp. 1024–1034). https://proceedings.neurips.cc/paper_files/paper/2017/file/8ef8326f3aa51aa5bfa8f0d0d2f37d3b-Paper.pdf
6. McKinney, W. (2010). *Data structures for statistical computing in Python*. In *Proceedings of the 9th Python in Science Conference* (pp. 51–56). <https://doi.org/10.25080/Majora-92bf1922-00a>
7. Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., ... & Duchesnay, E. (2011). *Scikit-learn: Machine learning in Python*. *Journal of Machine Learning Research*, 12, 2825–2830. <http://jmlr.org/papers/v12/pedregosa11a.html>

Appendices

Appendix A: Key Code Snippets

A1. Dataset Loading and Preprocessing

```
from torch_geometric.datasets import Planetoid
from torch_geometric.transforms import NormalizeFeatures

dataset = Planetoid(root='./data', name='Cora',
transform=NormalizeFeatures())
data = dataset[0]
```

A2. Train/Validation/Test Split

```
def create_masks(data, val_ratio=0.1, test_ratio=0.2):
    num_nodes = data.num_nodes
    indices = torch.randperm(num_nodes)

    test_size = int(num_nodes * test_ratio)
    val_size = int(num_nodes * val_ratio)
    train_size = num_nodes - val_size - test_size

    data.train_mask = torch.zeros(num_nodes, dtype=torch.bool)
    data.val_mask = torch.zeros(num_nodes, dtype=torch.bool)
    data.test_mask = torch.zeros(num_nodes, dtype=torch.bool)

    data.train_mask[indices[:train_size]] = True
    data.val_mask[indices[train_size:train_size + val_size]] = True
    data.test_mask[indices[train_size + val_size:]] = True

    return data

data = create_masks(data)
```

A3. GCN Model Definition

```
class GCN(nn.Module):
    def __init__(self, in_channels, hidden1, hidden2, out_channels):
        super(GCN, self).__init__()
        self.conv1 = GCNConv(in_channels, hidden1)
        self.bn1 = nn.BatchNorm1d(hidden1)
        self.conv2 = GCNConv(hidden1, hidden2)
        self.bn2 = nn.BatchNorm1d(hidden2)
        self.conv3 = GCNConv(hidden2, out_channels)
```

```

        self.skip = nn.Linear(in_channels, out_channels) if in_channels != out_channels else None
        self.dropout1 = nn.Dropout(0.5)
        self.dropout2 = nn.Dropout(0.5)
        self.act1 = nn.ELU()
        self.act2 = nn.ELU()

```

Appendix B: Training and Evaluation Loops

B1. Training Function

```

def train():
    model.train()
    optimizer.zero_grad()
    out = model(data.x, data.edge_index)
    loss = F.nll_loss(out[data.train_mask], data.y[data.train_mask])
    loss.backward()
    torch.nn.utils.clip_grad_norm_(model.parameters(), 2.0)
    optimizer.step()
    return loss.item()

```

B2. Evaluation Function

```

@torch.no_grad()
def evaluate(mask):
    model.eval()
    out = model(data.x, data.edge_index)
    pred = out[mask].argmax(dim=1)
    correct = pred.eq(data.y[mask]).sum().item()
    acc = correct / mask.sum().item()

    # Class-wise accuracy
    classes = torch.unique(data.y[mask])
    class_acc = {}
    for c in classes:
        class_mask = mask & (data.y == c)
        if class_mask.sum() > 0:
            class_correct =
pred[class_mask[mask]].eq(data.y[class_mask]).sum().item()
            class_acc[int(c)] = class_correct / class_mask.sum().item()

    return acc, class_acc

```

Appendix C: Sample Output Logs

C1. Final Accuracy Output Example

🎯 Final Test Accuracy: 0.8429

Class-wise Accuracy on Test Set:

Class Accuracy (%)

0	Class 0: 0.7500
1	Class 1: 0.7619
2	Class 2: 0.9070
3	Class 3: 0.8438
4	Class 4: 0.8438
5	Class 5: 0.9219
6	Class 6: 0.7879

C2. Early Stopping Message

Early stopping at epoch 2261!